



**FACULTY
OF MATHEMATICS
AND PHYSICS**
Charles University

BACHELOR THESIS

Daniel Crha

Board game with artificial intelligence

Department of Theoretical Computer Science and Mathematical Logic

Supervisor of the bachelor thesis: Mgr. Martin Pilát, Ph.D.

Study programme: Computer Science

Study branch: IOI

Prague 2020

I declare that I carried out this bachelor thesis independently, and only with the cited sources, literature and other professional sources. It has not been used to obtain another or the same degree.

I understand that my work relates to the rights and obligations under the Act No. 121/2000 Sb., the Copyright Act, as amended, in particular the fact that the Charles University has the right to conclude a license agreement on the use of this work as a school work pursuant to Section 60 subsection 1 of the Copyright Act.

In date
Author's signature

Most of all I want to thank my supervisor for his help and advice, he was always there for me whenever I needed his opinion. I also thank all of my family and friends for being there for me along the way, my journey has been long and I could not have done it without them.

Title: Board game with artificial intelligence

Author: Daniel Crha

Department: Department of Theoretical Computer Science and Mathematical Logic

Supervisor: Mgr. Martin Pilát, Ph.D., Department of Theoretical Computer Science and Mathematical Logic

Abstract: Multiplayer board games with imperfect information present a difficult challenge for many common game-playing algorithms. Studying their behavior in such games can be difficult, because existing implementations of such games have poor support for artificial intelligence. This thesis aims to implement an imperfect information multiplayer board game in a way that provides a framework for developing and testing artificial intelligences for board games with the aforementioned qualities. Furthermore, this thesis explores the implementation of several algorithms for the game. This aims to showcase the artificial intelligence framework, as well as to analyze the performance of existing algorithms when applied to a board game with elements such as hidden information and multiple players.

Keywords: board game, artificial intelligence

Contents

Introduction	3
Foreword	3
Goals	3
1 Related Work	4
1.1 Games and Frameworks	4
1.1.1 OpenAI Gym	4
1.1.2 boardgame.io	5
1.2 Algorithms	5
1.2.1 Minimax	5
1.2.2 Monte Carlo Methods	5
2 Game Design	6
2.1 High Level Design	6
2.2 Game Rules	7
2.2.1 Colonist Pick	7
2.2.2 Proper Turns	8
2.2.3 Colonists	8
2.2.4 Modules	9
3 Framework	10
3.1 Design	10
3.2 Interface	10
4 Used Algorithms	11
4.1 Random Decisions	11
4.2 Heuristics	11
4.3 MaxN	11
4.4 Information Set Monte Carlo Tree Search	11
5 Experiment Description	12
5.1 Game Balance Experiments	12
5.1.1 Description	13
5.1.2 Findings	13
5.2 Algorithm Comparison Experiments	15
5.2.1 Description	15
5.2.2 Findings	16
Conclusion	17
Bibliography	18
List of Figures	19
List of Tables	20

A	Attachments	21
A.1	User Documentation	21
A.1.1	Installation	21
A.1.2	User Interface	21
A.2	Developer Documentation	22
A.2.1	Prerequisites	22
A.2.2	Project Structure	22
A.2.3	Game Engine	22
A.2.4	Artificial Intelligence	22
A.2.5	User Interface	22
A.2.6	Experiments	22

Introduction

Foreword

In game theory, perfect information two-player games are often studied, and numerous algorithms have been designed with the purpose of playing them. This includes games like Chess and Go, which have had large breakthroughs in recent years [1]. However, real world situations do not always have perfect information, or only two parties involved. We could for example imagine multiple countries, which have only approximate information about the armies of their opponents. In this scenario, it could be useful to have tools to simulate potential enemy troop movements or placements.

Even though algorithms which are able to model imperfect information and multiple players are often useful, they are not studied nearly as often. Designing such an algorithm is not easy, and there are many pitfalls which make conventional game theory algorithms much less effective at solving imperfect information and multi-player problems. This thesis therefore aims to analyze the problems of implementing such algorithms, and to implement some of them in pursuit of that goal.

Naturally, some frameworks do already exist for the implementation of such games. However, at the time of writing, some of them only have AI (Artificial Intelligence) support as an experimental and sparsely documented feature [2], and others only focus on specific fields of AI [3]. This work aims to provide a kind of “plug-and-play” experience, where AI developers have minimal barriers between cloning a git repository and having a working AI.

Goals

The main goal of this thesis is to create a multi-player board game with imperfect information states. The game’s name is *Colonizers*. The game will primarily be designed with AI in mind, and it will provide a reasonable interface for the implementation of AI players.

Another goal is the implementation of several AI players for said game. This will allow us to not only explore potential problems with implementing AIs for games of this kind. We will also verify that the API (Application Programming Interface) provided by the game is sufficient for implementation of such AI players, and that the API is reasonably easy to use.

1. Related Work

Before we discuss the design of *Colonizers* and its implementation, let us first make an overview of existing related work. We are interested mainly in two areas here:

- Implementation of similar games, and frameworks facilitating that
- Algorithms adapted for multi-player games and algorithms adapted for imperfect information games

1.1 Games and Frameworks

1.1.1 OpenAI Gym

OpenAI Gym is "a toolkit for developing and comparing reinforcement learning algorithms" [3]. It is a popular tool in the reinforcement learning field, because it is modular, and easy to work with. It features a standardized API for all of its environments (games or problems). This means that agents built for one environment can be easily transitioned to other environments, without having to structurally rebuild it. Another benefit is the fact that it is easy to create new environments, and these newly created environments can be used by anyone, since the API is standardized.

The following is an example (as presented in the OpenAI Gym documentation [3]) of a Python program which solves one of the simpler environments available out-of-the-box in OpenAI Gym:

```
import gym
env = gym.make('CartPole-v0')
for i_episode in range(20):
    observation = env.reset()
    for t in range(100):
        env.render()
        print(observation)
        action = env.action_space.sample()
        observation, reward, done, info = env.step(action)
        if done:
            print("Episode finished after {} timesteps".format(t+1))
            break
    env.close()
```

The environment being solved (*CartPole-v0*) is a task where the AI must balance a pole by moving the cart below it left and right. The agent only performs random moves, but the example clearly illustrates how the agent interacts with the environment.

OpenAI Gym is not particularly suitable for the study of multi-player games with imperfect information for a few reasons:

- It only supports reinforcement learning agents. The API is designed with this in mind, and does not provide support for any other machine learning methods.

- It does not provide any tools for determinization¹ of imperfect information states. This would force AIs to track their own information sets, and to then produce determinizations of game states on their own.

In spite of that, there is something we can take away from OpenAI Gym when designing *Colonizers*. Notably, the API is very elegant, and creating an AI which simply plays random moves is a matter of very few lines of code. We will try to achieve this with *Colonizers*.

1.1.2 boardgame.io

boardgame.io [2] is a game engine for creating turn-based games. It features many helpful features for creating board games, such as support for multiplayer, randomness, imperfect information, and a few other useful features.

Using boardgame.io for the implementation of *Colonizers* would make many things much simpler, notably the implementation of game logic would be trivial. However, it is also not suitable for the purposes of this thesis, because the AI support is poor. The engine does feature a degree of AI support, but the API is limited to using pre-existing AIs which ship with the game (an MCTS AI and a random AI). The API only provides a method for us to list the legal moves in a given game state — it does not however provide ways to implement a fully custom AI.

1.2 Algorithms

1.2.1 Minimax

1.2.2 Monte Carlo Methods

¹By determinization of a game state, we understand the conversion of a game state with hidden information into a game state with perfect information. Determinization takes into account the information set of the given player. For example, we can imagine a poker player who has been dealt a hand which includes the Queen of Hearts. When this player is thinking about what other players may have, the Queen of Hearts is out of the question, since the player has it, and there is only one in the deck. Therefore, a rational determinization of a poker game state would be to take all cards which started in the deck, remove the ones the player is holding, and then randomly assign other cards to the other players.

2. Game Design

This chapter’s purpose is to discuss the considerations which went into designing the game’s rules, and to describe said rules in detail.

The game is set on Mars with futuristic themes. In the game universe, Mars is only just starting to be settled by humans, and there was a precious mineral found under the surface - Omnium. This triggered a rush of colonists, who are eager to make some profit. In the game, they compete for resources, and they all want to build the largest colony, because the person with the largest colony can extract the most Omnium and get rich.

2.1 High Level Design

Colonizers has a few design decisions which are inherently set in stone by the premise of this thesis:

- The game must have more than two players
- The game must feature hidden information
- The game must support AI

AI support is only tangentially related to the design of the game rules, therefore we will not discuss it at length in this chapter. We will focus on the other two requirements.

Colonizers is a four-player game. Four was chosen as a sweet spot for complexity, since with five players, the game would start to get prohibitively expensive to compute. It could be argued that three would accomplish the same goal, but four makes more sense with respect to having enough design space. Four players is also a very common player count for board games.

Hidden information is somewhat more tricky to get right. It can be implemented in many ways, but even the simplest inclusions make the game much more tricky to process with AI. There are two elements of hidden information in *Colonizers*:

- Players’ hands and the Deck
- Players’ colonists

These elements will be explained in more detail in the following section. It is worth noting however that it is possible to make information-gathering plays, even though only in very limited ways and in rare circumstances. Information-gathering plays are not a large design focus of *Colonizers*.

The game also features interaction between players — both malicious and cooperative. This naturally means that it is possible for multiple players to cooperate in order to gain an advantage, or to conspire against another player in order to damage that player’s chances of winning. Many traditional AI algorithms are not capable of cooperation or conspiracy, which provides room for specialized AIs to shine.

2.2 Game Rules

The game is played in rounds, which comprise of turns. Turns then comprise of phases. The four players start the game in a given order, and they always take turns in this order for the entire game.

Each player has a colony where they can build modules (the game’s terminology for buildings). Each module has a value, and the goal of the game is to build the most valuable colony. When any player builds eight modules in their colony, the game ends at the end of that round. When the game ends, the values of all modules in each player’s colony are added up, and the resulting value is that player’s final score. Players can also get a bonus to their score if they reached eight buildings in their colony before the game ended — the first player to build eight modules gets four bonus points, and other players to build eight modules get two bonus points each. The final ranking of the players when the game ends is determined by points — players with more points rank higher. If multiple players are tied in points, the player whose position according to the player order is earlier ranks higher. More information about modules and ways to interact with them follow in subsequent subsections.

There is also a rare game end condition, which is triggered by players attempting to draw from an empty deck. This immediately ends the game with a draw, giving all players zero points and a rank of zero.

2.2.1 Colonist Pick

At the start of each round, players take turns picking colonists. A colonist is a character with special powers, and the player controls a given colonist for one turn. A player’s colonist is hidden from the other players. There are six colonists in the game (see subsection 2.2.3 for a list of available colonists). At the start of the colonist picks, a random colonist is secretly removed for play for the rest of the round. Then, players take turns picking from the remaining colonists one by one. This means that after the last player picks, there is one colonist left over. This colonist is then removed from play for the rest of the round.

The colonist pick phase creates a situation where players have asymmetrical information. For example, the first player knows which colonist was removed at the start, but has no information about the other players apart from knowing the four colonist he is passing on. In contrast, the last player has relatively little information about the players before them, but they know which colonist is removed from play after being left over.

2.2.2 Proper Turns

After each player has chosen a colonist, the players take their actual turns, in order of first to last. Each player acts in all phases of their turn before passing the turn to the next player.

Each turn is comprised of the following phases:

- *Draw Phase.* The player may choose between acquiring two Omnium (the game's currency) and drawing 2 modules from the deck. The player must then keep one of the modules, and place the other at the bottom of the deck. The drawing action is not available if the player's hand is full (five modules). The player's colonist is also revealed to other players during this phase.
- *Power Phase.* The player may choose to use their colonist's active ability if the colonist has one.
- *Build Phase.* The player may choose to build one module from their hand. To build a module, the player must spend the Omnium amount required by the module's build cost. Building the module adds it to the player's colony and removes it from their hand.

2.2.3 Colonists

The following colonists and their respective abilities are available in *Colonizers*:

- Visionary
 - Passive Ability: Draw a card if the player's hand is not full (maximum hand capacity is five).
 - Active Ability: None
- Ecologist
 - Passive Ability: Gain 1 Omnium for each green module in his colony.
 - Active Ability: None
- Miner
 - Passive Ability: Gain 1 Omnium for each blue module in his colony.
 - Active Ability: None
- General
 - Passive Ability: Gain 1 Omnium for each red module in his colony.
 - Active Ability: None
- Opportunist
 - Passive Ability: None
 - Active Ability: Steal up to 2 Omnium from a chosen colonist. If no player controls the chosen colonist, this ability has no effect.

- Spy
 - Passive Ability: None
 - Active Ability: Swap hands with a chosen colonist. If no player controls the chosen colonist, this ability has no effect.

2.2.4 Modules

The deck starts with 52 modules. The following is a table of all modules available:

Name	Build Cost	Value	Color	Quantity
Oxygen Generator	4	4	Green	4
Water Reservoir	5	6	Green	4
Hydroponics Facility	6	8	Green	4
Eco-Dome	8	11	Green	1
Marketplace	2	2	Blue	4
Warehouse	3	3	Blue	4
Quarry	5	6	Blue	4
Omnium Purification Plant	8	10	Blue	1
Garrison	1	1	Red	4
Barracks	2	2	Red	4
Military Academy	3	3	Red	4
Planetary Defense System	6	7	Red	1
Housing Unit	1	1	None	4
Spaceport	4	5	None	4
Research Lab	6	8	None	4
Mass Relay	12	16	None	1

Table 2.1: Available modules.

3. Framework

3.1 Design

3.2 Interface

4. Used Algorithms

4.1 Random Decisions

4.2 Heuristics

4.3 MaxN

4.4 Information Set Monte Carlo Tree Search

5. Experiment Description

There are two qualities which we want to analyze with respect to the game and the implemented AI algorithms:

- Identify potential asymmetries in game balance
- Compare methodologies used by the AI algorithms

To this end, we conducted five experiments, split according to their purpose. The following sections elaborate on the experiments and their results.

5.1 Game Balance Experiments

As mentioned in chapter 2, the game features a degree of asymmetry. The order in which players take their turns inherently changes the viability of certain strategies, because players in different positions have different information sections available to them. For example, the player in the first position has perfect information about which colonist was removed from play during the colonist pick phase, while the second and third players do not have such certainty.

Most importantly however, a player's colonist is revealed at the start of their turn. This means that if the player in the fourth position is a Spy or an Opportunist, they will know all the other players' colonists when their turn comes around. This means that this player will be able to target any player with their targeted ability without the fear of missing or hitting an unintended target.

With these things in mind, we can hypothesize that players in the earlier positions have an easier time achieving synergy-based strategies, since they get priority when picking colonists. On the other hand, we can also hypothesize that players in later positions will benefit from play based around using targeted colonist abilities.

In Chess, it is widely agreed that the white player has an advantage [4]. Similarly, we aim to discover whether player ordering confers a measurable advantage to any player in *Colonizers*. We will conduct this experiment with the null hypothesis — we assume that there is no significant advantage for any player ordering.

5.1.1 Description

We conducted two experiments in this section. In both of them, four identical AIs played 1000 games against each other. In the first experiment, the AI in question was `RandomIntelligence`, and in the second experiment it was `HeuristicIntelligence`.

All random events were seeded, and the results of the games were captured in JSON (JavaScript Object Notation) files. The results were then parsed and analyzed. The JSON result files can be found in the attached source code, refer to subsection A.2.6 for more information on their location and semantics.

The random seeds used by application components during the experiment were as follows:

- `RandomIntelligence`: seed 42
- `HeuristicIntelligence`: seed 97
- `GameConstants`: seed was changed every game to prevent the same game from being played 1000 times. The seeds were generated by a C# random number generator seeded with 42.

5.1.2 Findings

Experiment 1

First off, let us focus on the experiment runs with `RandomIntelligence`. Results of the 1000 runs can be seen in table 5.1.

Position	1	2	3	4
Wins	310	213	251	226
Losses	197	261	279	263
Average rank	2.3	2.572	2.553	2.575

Table 5.1: Results of `RandomIntelligence` play.

The most notable result we have is the fact that AIs in the first position seem to be winning the most often. AIs in the first position also lose (place fourth) less, and they have a better average ranking overall.

We can try to verify the significance of these results mathematically. If we assume that the rank at the end of the game follows a normal distribution, we can compute a confidence interval. Let \bar{x} be the sample mean, let s be the sample standard deviation and let n be the sample size. We are looking for a confidence interval for the unknown mean μ . The $100(1 - \alpha)\%$ confidence interval for μ is

$$\hat{\mu}_L = \bar{x} - z_{\alpha/2} \cdot s / \sqrt{n}, \quad \hat{\mu}_U = \bar{x} + z_{\alpha/2} \cdot s / \sqrt{n},$$

The sample mean for rank among first position AIs is 2.3 as seen in table 5.1, and the sample standard deviation is 1.1069. If we want a 95% confidence interval, we will use $z_{\alpha/2} = 1.96$. This gives us the confidence interval of

$$\hat{\mu}_L = 2.2314, \quad \hat{\mu}_U = 2.3686$$

We can also compute a 95% confidence interval for the mean rank of AIs in position 3:

$$\hat{\mu}_L = 2.4821, \hat{\mu}_U = 2.6239$$

These intervals do not overlap, therefore we can conclude that there is a statistically significant difference between the rank means among AIs at different positions. This may indicate a potential balance issue in the rules of the game, with the first position being more powerful than the other ones, which are similar in power. However, measurement on randomly choosing AIs does not necessarily indicate imbalance, since random agents do not play optimal strategies. Therefore we cannot conclude anything about game balance just yet, but this statistical difference is worth keeping in mind.

Experiment 2

The other experiment in this section is very similar to the first one, except we have four instances of `HeuristicIntelligence` instead of four instances of `RandomIntelligence` playing against each other. Results of the 1000 runs can be seen in table 5.2.

Position	1	2	3	4
Wins	230	202	282	286
Losses	415	298	152	135
Average rank	2.8	2.67	2.302	2.228

Table 5.2: Results of `HeuristicIntelligence` play.

If we compare these results to those in table 5.1, we can see almost exactly the opposite results. With random AIs playing, we saw that the AI in the first position had a statistically significant advantage. With heuristically driven AIs, it is obvious on first glance that earlier positions are less powerful and later positions are more powerful. We can verify this statistically by computing confidence intervals for the first and fourth ranks. The 99% confidence interval for the first position is

$$\hat{\mu}_L = 2.7019, \hat{\mu}_U = 2.8981$$

while the 99% confidence interval for the fourth position is

$$\hat{\mu}_L = 2.1458, \hat{\mu}_U = 2.3102$$

The intervals do not overlap, therefore we can conclude that there is a statistically significant difference between the means of these positions' respective average ranks.

While the differences between wins per position are notable, the most interesting are the loss statistics. It would appear that the earlier ranks (particularly the first one) are susceptible to being targeted by players in other ranks. Since every player's colonist is revealed at the start of their turn, this makes the first position an easy target for all other players. The game does have counter-balances for this situation — notably the fact that the first player to build their colony to full gets four extra victory points. However, it would seem that the heuristic AI does not

have the necessary tools to deal with being targeted down by others. This could possibly be due to an implementation bias inherent in the chosen heuristics, but it could also signal a game balance issue.

5.2 Algorithm Comparison Experiments

We have implemented four algorithms in this thesis — `RandomIntelligence`, `HeuristicIntelligence`, `MaxnIntelligence` and `ISMCTSIntelligence`. In order to determine the qualities of said algorithms, we will analyze their differences, along with their advantages and disadvantages. We will also look at how the algorithms perform in play against each other, with the hopes of determining which algorithm is the most suitable for a game like *Colonizers*.

To start with, we would not expect `RandomIntelligence` to perform well in any kind of mutual play. It is present simply as a benchmark for the performance of other AIs.

The more important benchmark is `HeuristicIntelligence`, since it represents rules which were created by observing humans play the game. Therefore we consider this AI to be a minimum benchmark for other AIs to be competent.

`MaxnIntelligence` is based on the MaxN algorithm [5], which is itself based on Minimax [6]. This AI was adapted for imperfect information games, and it spends a non-trivial amount of computing power on simply exploring possible determinizations of the current game state. If we take that into consideration, along with the fact that the branching factor for *Colonizers* is non-trivial, we would expect `MaxnIntelligence` to perform relatively poorly. The depth of the search trees used could not be reasonably increased beyond 7, due to performance concerns. We expect that any kind of long-term strategy could not be achieved by it since it lacks the necessary exploration depth. The Minimax family of algorithms does however offer very solid insight into the few turns it examines, therefore we hypothesize that this AI will be primarily good at tactics-based play. The performance of `MaxnIntelligence` against `HeuristicIntelligence` is uncertain, therefore we will follow the null hypothesis and assume that their performances are statistically similar. We also hypothesize that since this AI has a strong foundation for tactical prowess, it should win more often when in later positions (namely third and fourth).

The final AI tested is `ISMCTSIntelligence`. This AI is well-adapted to imperfect information and multiple player environments. Therefore we would expect it to outperform the three aforementioned AIs in most situations. We expect it to play well in most circumstances, regardless of player permutation.

5.2.1 Description

We conducted three experiments in this section. In the first experiment, one of each implemented intelligence played 50 games against each other. This experiment is meant to assess the general playing ability of the AIs. In the second experiment, we let two instances of `HeuristicIntelligence` and two instances of `MaxnIntelligence` play against each other for 50 games. Lastly in the third experiment, we performed the same thing as in the second experiment, but we replaced `MaxnIntelligence` instances with `ISMCTSIntelligence` instances. These

two experiments are meant to benchmark the adapted algorithms against the heuristic solution.

All random events were seeded, and the results of the games were captured in JSON files. The results were then parsed and analyzed. The JSON result files can be found in the attached source code, refer to subsection A.2.6 for more information on their location and semantics.

The random seeds used by application components during the experiment were as follows:

- `RandomIntelligence`: seed 42
- `HeuristicIntelligence`: seed 97
- `MaxnIntelligence`: seed 99
- `ISMCTSIntelligence`: seed 15
- `GameConstants`: seed was changed every game to prevent the same game from being played 1000 times. The seeds were generated by a C# random number generator seeded with 42.

5.2.2 Findings

WORK IN PROGRESS

Experiment 3

Experiment 4

Experiment 5

Conclusion

Bibliography

- [1] David Silver, Aja Huang, Christopher Maddison, Arthur Guez, Laurent Sifre, George Driessche, Julian Schrittwieser, Ioannis Antonoglou, Veda Panneershelvam, Marc Lanctot, Sander Dieleman, Dominik Grewe, John Nham, Nal Kalchbrenner, Ilya Sutskever, Timothy Lillicrap, Madeleine Leach, Koray Kavukcuoglu, Thore Graepel, and Demis Hassabis. Mastering the game of go with deep neural networks and tree search. *Nature*, 529:484–489, 01 2016.
- [2] boardgame.io [online]. <https://boardgame.io>. Accessed: 2020-05-31.
- [3] Openai gym [online]. <https://gym.openai.com>. Accessed: 2020-05-31.
- [4] W. F. Streeter. Is the first move an advantage? *Chess Review*, page 16, 05 1946.
- [5] C. A. Luckhardt and K. B. Irani. An algorithmic solution of n-person games. In *AAAI*, 1986.
- [6] I. Millington and J. Funge. *Artificial Intelligence for Games*. Second Edition. Morgan Kaufmann, 2009.

List of Figures

List of Tables

2.1	Available modules.	9
5.1	Results of <code>RandomIntelligence</code> play.	13
5.2	Results of <code>HeuristicIntelligence</code> play.	14

A. Attachments

A.1 User Documentation

A.1.1 Installation

A.1.2 User Interface

A.2 Developer Documentation

A.2.1 Prerequisites

A.2.2 Project Structure

A.2.3 Game Engine

A.2.4 Artificial Intelligence

A.2.5 User Interface

A.2.6 Experiments